



university of  
 groningen

faculty of science  
 and engineering

bernoulli institute


# STATEFUL DATA ANALYTICS OVER PROGRAMMING MODELS OF NETWORKS

Author : Gheorghe Pojoga


Supervisors : Prof. Dr. Boris Koldehofe, Bochra Boughzala




# Networking - a way of life




# Networking - a way of life




# Networking - a way of life





# Networking - a way of life





# Networking - a way of life



# Networking - a way of life





# Networking - a way of life







# Networking - a way of life




# Networking - a way of life



# Networking - a way of life



# Networking - a way of life



# Programmable Networks !

# Programmable Networks !



# Programmable Networks !



# Programmable Networks !






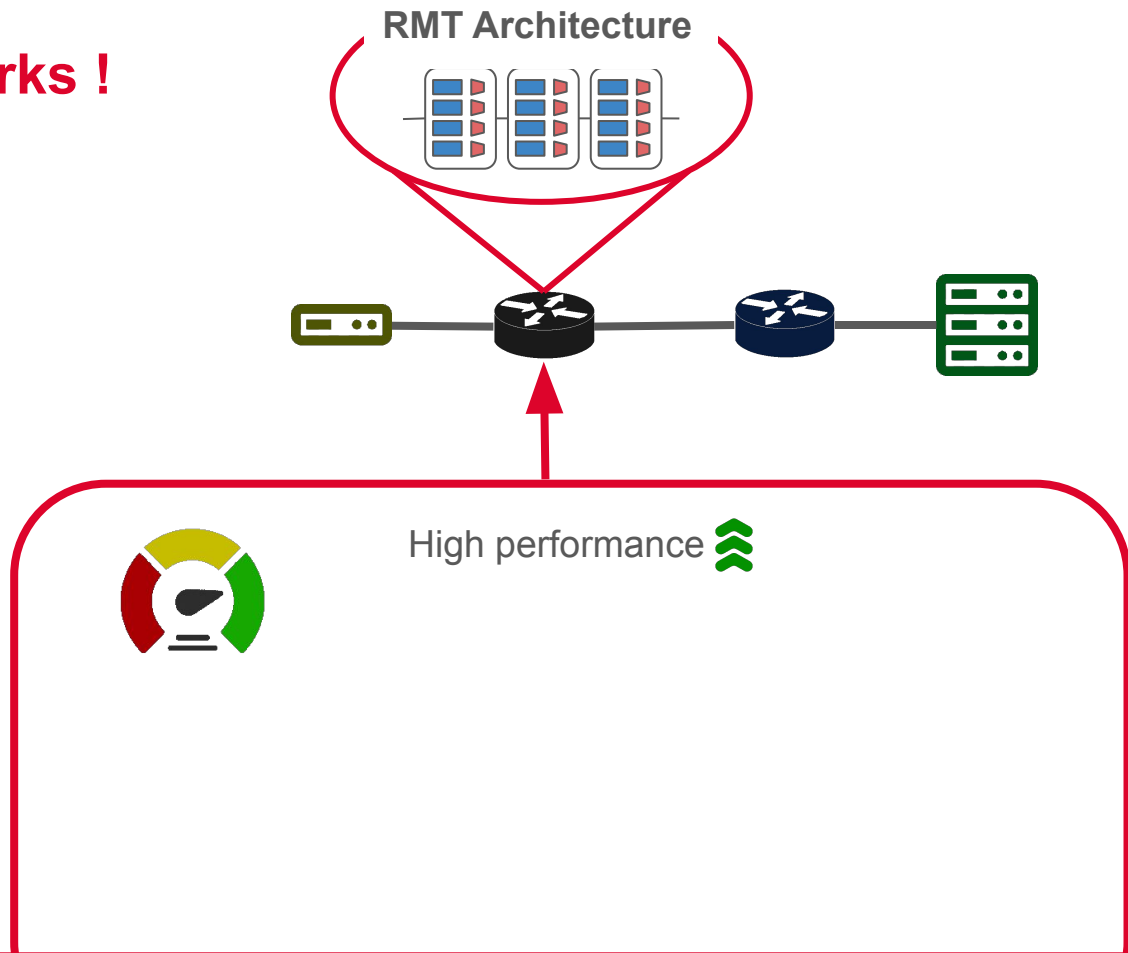
# Programmable Networks !



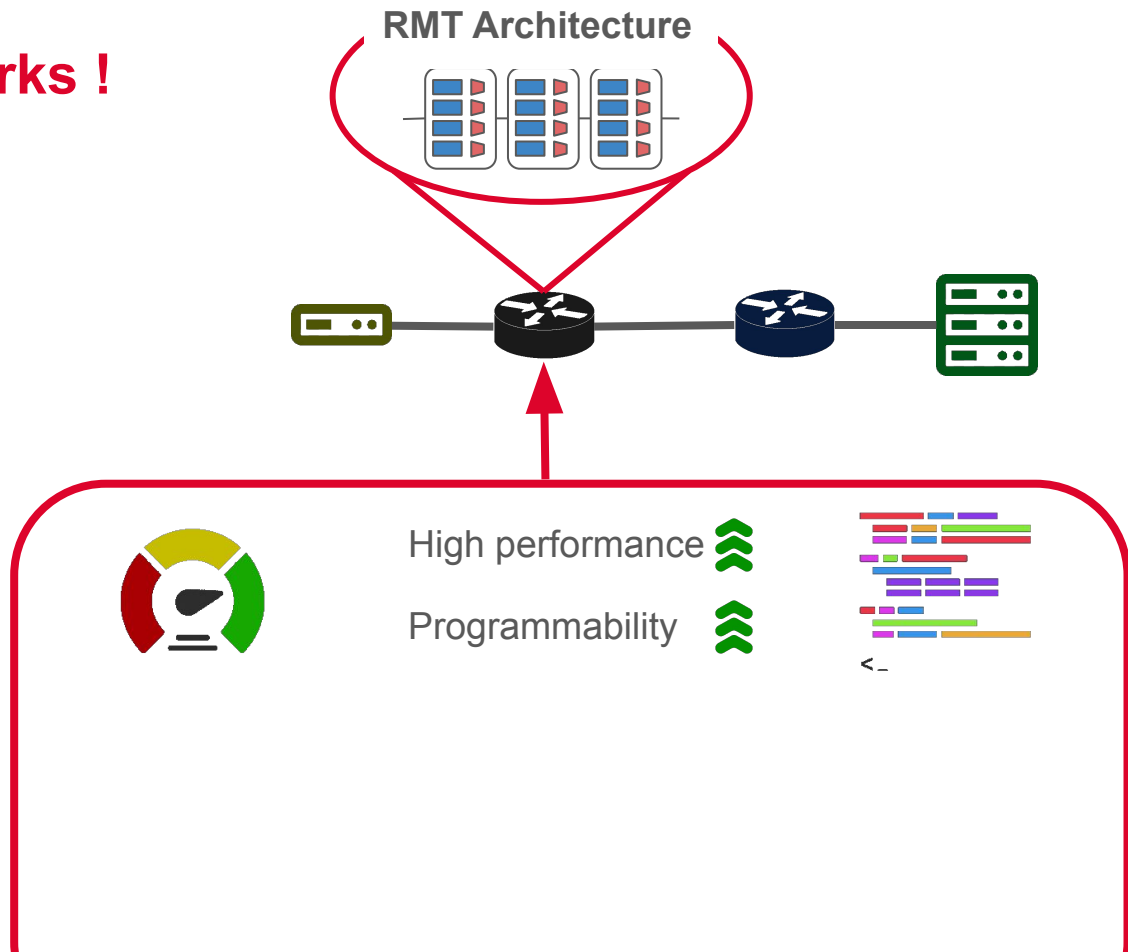
# Programmable Networks !



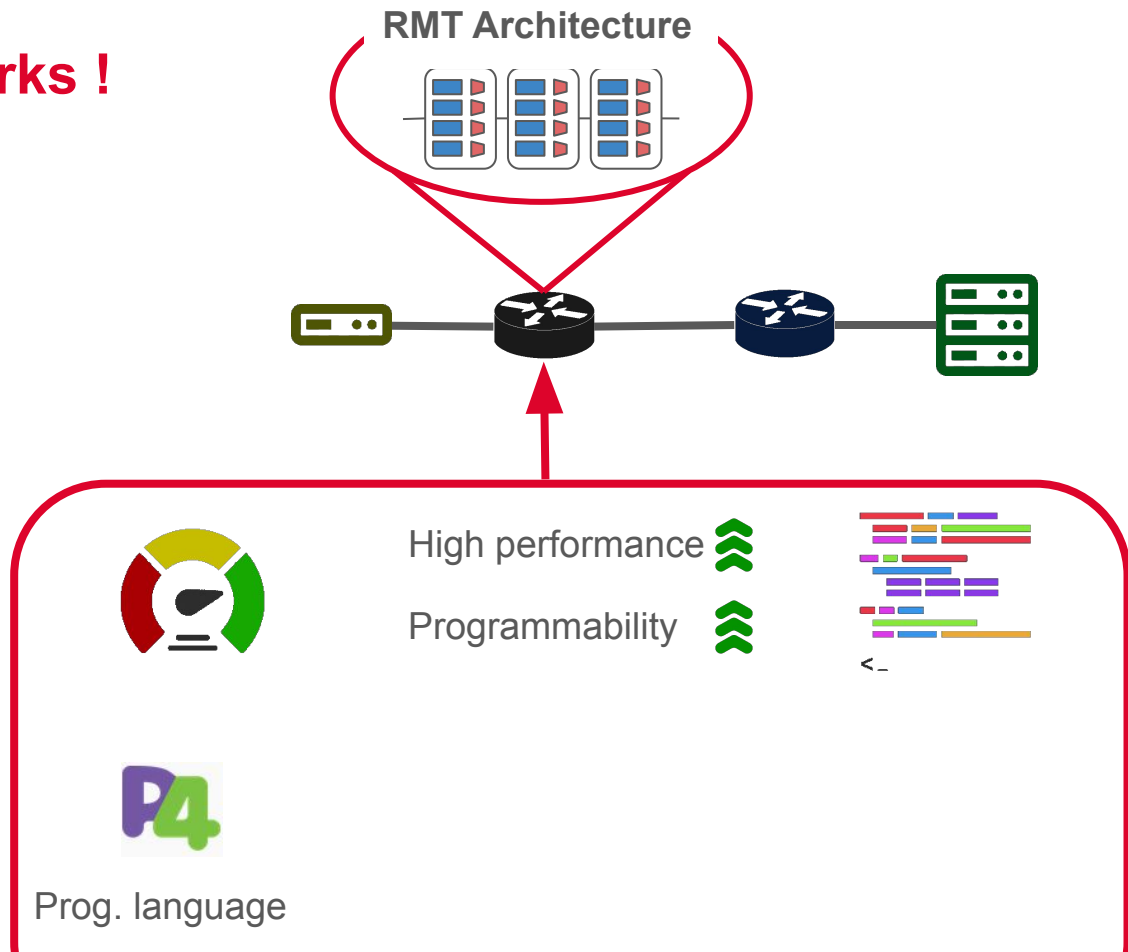
# Programmable Networks !



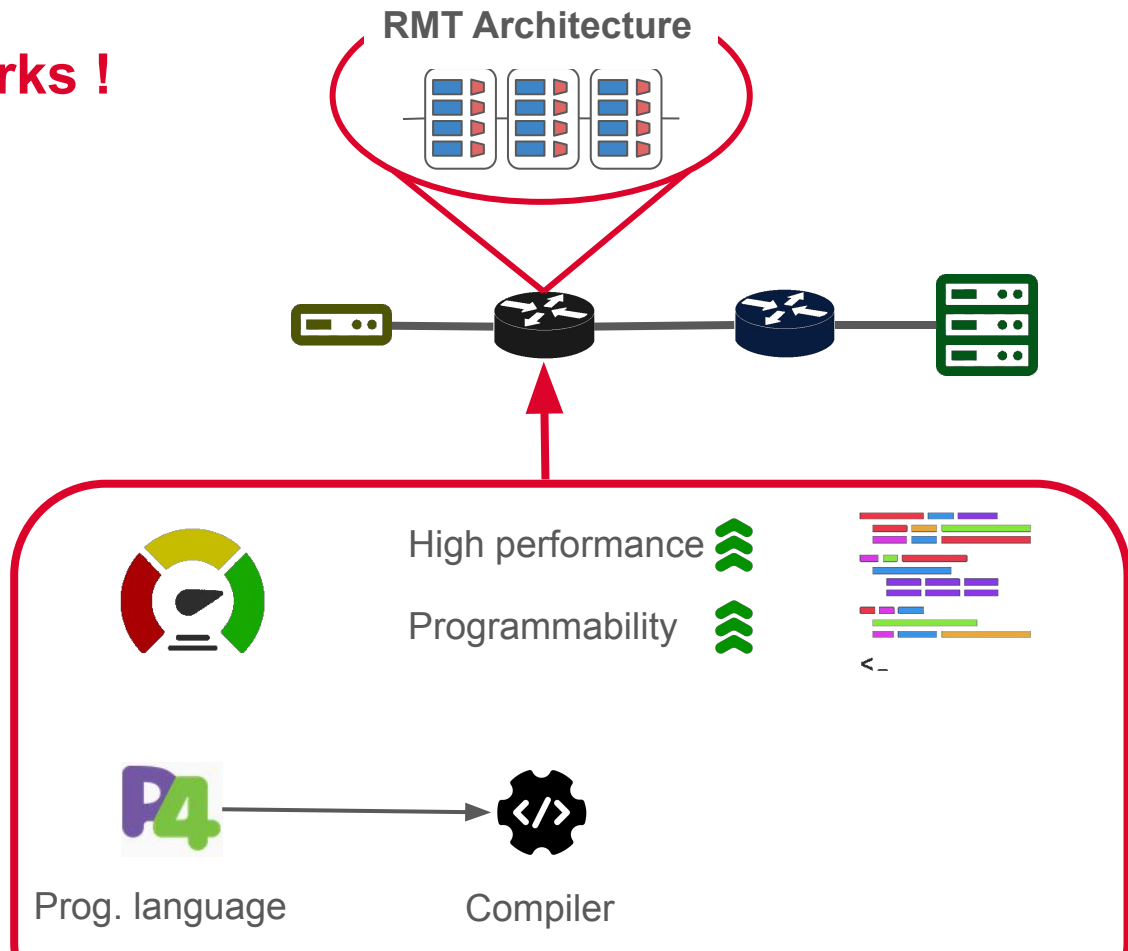
# Programmable Networks !



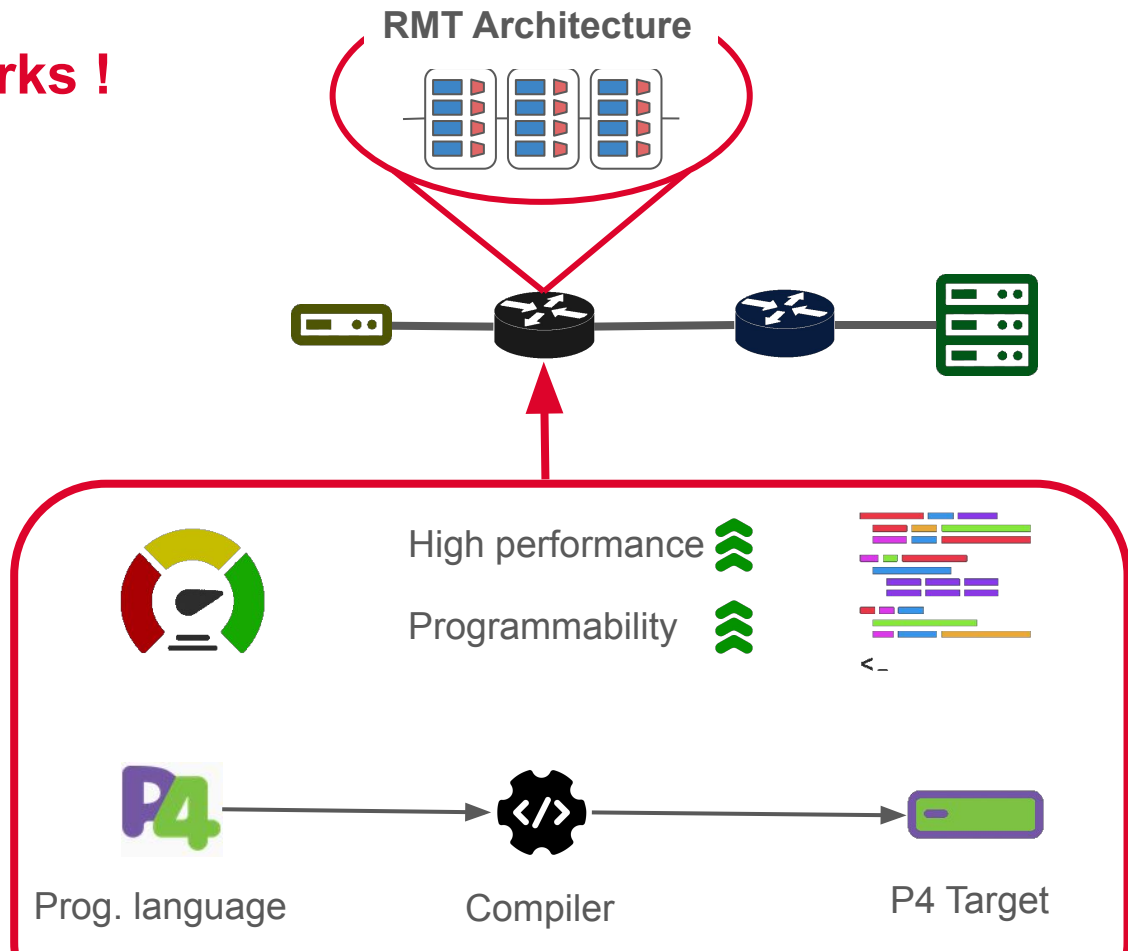
# Programmable Networks !



# Programmable Networks !




# Programmable Networks !



# Programmable Networks !

## Software Defined Networking



High performance >>>

Programmability >>>



Prog. language



Compiler



P4 Target




# Programmable Networks !

Software Defined Networking

Data Plane

Control Plane

RMT Architecture



High performance



Programmability



Prog. language



Compiler



P4 Target



university of  
groningen

faculty of science  
and engineering

bernoulli institute

# Programmable Networks !


## Software Defined Networking

Data Plane

Control Plane




## RMT Architecture



High performance

Programmability



Prog. language



Compiler



P4 Target


# Programmable Networks !

## Software Defined Networking


### Data Plane



### Control Plane



## RMT Architecture




High performance



Programmability



Prog. language



Compiler



P4 Target

# P4 Constraints

# P4 Constraints



# P4 Constraints




- ✗ Floating Point Numbers
- ✗ Loops

# P4 Constraints



✗ Floating Point Numbers

✗ Loops



# P4 Constraints




Java™



Floating Point Numbers



Loops






# P4 Constraints



✗ Floating Point Numbers

✗ Loops



# P4 Constraints




Java™



Floating Point Numbers



Loops



# P4 Constraints




Java™



Floating Point Numbers



Loops




# P4 Constraints



Floating Point Numbers



Loops




# P4 Constraints



✗ Floating Point Numbers

✗ Loops




# P4 Constraints



✗ Floating Point Numbers

✗ Loops




# P4 Constraints



Floating Point Numbers



Loops



## Limitations :

- Constant time complexity is required
- Limited amount of memory

# Research Questions



# Research Questions



**What are the acceleration capabilities stateful data analytics inside the network ?**

# Research Questions



**What are the acceleration capabilities stateful data analytics inside the network ?**



**What are the limitation of the in-network computing model ?**

# Research Questions



What are the acceleration capabilities stateful data analytics inside the network ?



What are the limitation of the in-network computing model ?



Is it viable to offload stateful end-host operators onto network devices ?

# Sliding Window Average

# Sliding Window Average

...	2	1	7	4	5
-----	---	---	---	---	---

# Sliding Window Average

...	2	1	7	4	5
-----	---	---	---	---	---

...	1	7	4	5	3
-----	---	---	---	---	---

# Sliding Window Average

...	2	1	7	4	5
-----	---	---	---	---	---

...	1	7	4	5	3
-----	---	---	---	---	---

...	7	4	5	3	9
-----	---	---	---	---	---

# Sliding Window Average

...	2	1	7	4	5
-----	---	---	---	---	---

...	1	7	4	5	3
-----	---	---	---	---	---

...	7	4	5	3	9
-----	---	---	---	---	---

## P4-target Constraints

- ✗ Division at runtime
- ✗ Floating point numbers



# Sliding Window Average

...	2	1	7	4	5
-----	---	---	---	---	---

...	1	7	4	5	3
-----	---	---	---	---	---

...	7	4	5	3	9
-----	---	---	---	---	---

Sum = 17

Window size = 3

## P4-target Constraints

- ✗ Division at runtime
- ✗ Floating point numbers

# Sliding Window Average



Sum = Sum - 5

Sum = 17

Window size = 3

## P4-target Constraints

- ✗ Division at runtime
- ✗ Floating point numbers

# Sliding Window Average

...	2	1	7	4	5
-----	---	---	---	---	---

...	1	7	4	5	3
-----	---	---	---	---	---

...	7	4	5	3	9
-----	---	---	---	---	---

Sum = 17

Window size = 3

...	4	5	3	9	7
-----	---	---	---	---	---

Sum = Sum - 5

Sum = Sum + 7

## P4-target Constraints

- ✗ Division at runtime
- ✗ Floating point numbers

# Sliding Window Average

...	2	1	7	4	5
-----	---	---	---	---	---

...	1	7	4	5	3
-----	---	---	---	---	---

...	7	4	5	3	9
-----	---	---	---	---	---

Sum = 17

Window size = 3

...	4	5	3	9	7
-----	---	---	---	---	---

Sum = Sum - 5

Sum = Sum + 7

## P4-target Constraints

- ✗ Division at runtime
- ✗ Floating point numbers

Time Complexity :  $O(1)$

Memory Complexity :  $O(n)$ ,  $n$  - the size of the window

# Median

# Median

## Constraints:

- Only integers are accepted
- The set of possible values must be known in advance



# Median

## Constraints:

- Only integers are accepted
- The set of possible values must be known in advance

1	4	5	7	9
---	---	---	---	---

# Median

## Constraints:

- Only integers are accepted
- The set of possible values must be known in advance

1	4	5	7	9
---	---	---	---	---

If (val < median) : Move Left

1	4	5	7	9
---	---	---	---	---

If (val == median) : Stay

1	4	5	7	9
---	---	---	---	---

If (val > median) : Move Right

1	4	5	7	9
---	---	---	---	---



# Median

## Constraints:

- Only integers are accepted
- The set of possible values must be known in advance

1	4	5	7	9
---	---	---	---	---

If (val < median) : Move Left

1	4	5	7	9
---	---	---	---	---

If (val == median) : Stay

1	4	5	7	9
---	---	---	---	---

If (val > median) : Move Right

1	4	5	7	9
---	---	---	---	---

## P4-target Constraints

- ✗ Nontrivial operations
- ✗ Bit shift of more than 8 bits

# Median

## Constraints:

- Only integers are accepted
- The set of possible values must be known in advance

1	4	5	7	9
---	---	---	---	---

If (val < median) : Move Left

1	4	5	7	9
---	---	---	---	---

If (val == median) : Stay

1	4	5	7	9
---	---	---	---	---

If (val > median) : Move Right

1	4	5	7	9
---	---	---	---	---

Time Complexity :  $O(1)$

Memory Complexity :  $O(n)$ ,  $n$  - the number of possible values

## P4-target Constraints

- ✗ Nontrivial operations
- ✗ Bit shift of more than 8 bits

# Performance Comparison

# Performance Comparison

**In-Network Computing**

**End-Host Computing**



university of  
 groningen

faculty of science  
 and engineering

bernoulli institute


# Performance Comparison

In-Network Computing

End-Host Computing

Create / Remove a flow (Latency)

P4




# Performance Comparison

In-Network Computing

End-Host Computing

Create / Remove a flow (Latency)

$$P4 \quad \left( \sum_{i=1}^{n+1} L_i + \sum_{i=1}^n S_i \right)$$



# Performance Comparison


In-Network Computing

End-Host Computing



Create / Remove a flow (Latency)

$$\text{P4} \quad \left( \sum_{i=1}^{n+1} L_i + \sum_{i=1}^n S_i \right) + \left( \sum_{i=1}^{k+1} L_{ci} + \sum_{i=1}^k S_{ci} \right)$$



# Performance Comparison


In-Network Computing

End-Host Computing



Create / Remove a flow (Latency)

$$\text{P4} \quad \left( \sum_{i=1}^{n+1} L_i + \sum_{i=1}^n S_i \right) + 2 \left( \sum_{i=1}^{k+1} L_{ci} + \sum_{i=1}^k S_{ci} \right) + T_{\text{controller}} + T_{\text{P4OS}}$$





# Performance Comparison


In-Network Computing

End-Host Computing



Create / Remove a flow (Latency)

$$\text{P4} \quad 2 \left( \sum_{i=1}^{n+1} L_i + \sum_{i=1}^n S_i \right) + 3 \left( \sum_{i=1}^{k+1} L_{ci} + \sum_{i=1}^k S_{ci} \right) + T_{\text{controller}} + T_{\text{P4OS}}$$



# Performance Comparison


In-Network Computing

End-Host Computing




Create / Remove a flow (Latency)

$$\text{P4} \quad 2 \left( \sum_{i=1}^{n+1} L_i + \sum_{i=1}^n S_i \right) + 3 \left( \sum_{i=1}^{k+1} L_{ci} + \sum_{i=1}^k S_{ci} \right) + T_{\text{controller}} + T_{\text{P4OS}}$$



DPDK



# Performance Comparison

In-Network Computing



End-Host Computing



Create / Remove a flow (Latency)

$$\text{P4} \quad 2 \left( \sum_{i=1}^{n+1} L_i + \sum_{i=1}^n S_i \right) + 3 \left( \sum_{i=1}^{k+1} L_{ci} + \sum_{i=1}^k S_{ci} \right) + T_{\text{controller}} + T_{\text{P4OS}}$$

$$\text{DPDK} \quad \left( \sum_{i=1}^{n+1} L_i + \sum_{i=1}^n S_i \right)$$



# Performance Comparison

In-Network Computing



End-Host Computing



Create / Remove a flow (Latency)

$$\text{P4} \quad 2 \left( \sum_{i=1}^{n+1} L_i + \sum_{i=1}^n S_i \right) + 3 \left( \sum_{i=1}^{k+1} L_{ci} + \sum_{i=1}^k S_{ci} \right) + T_{\text{controller}} + T_{\text{P4OS}}$$

$$\text{DPDK} \quad \left( \sum_{i=1}^{n+1} L_i + \sum_{i=1}^n S_i \right) + \left( \sum_{i=1}^{m+1} L_{Si} + \sum_{i=1}^m S_{Si} \right) + T_{\text{DPDK}}$$



# Performance Comparison

In-Network Computing



End-Host Computing



Create / Remove a flow (Latency)

$$\text{P4} \quad 2 \left( \sum_{i=1}^{n+1} L_i + \sum_{i=1}^n S_i \right) + 3 \left( \sum_{i=1}^{k+1} L_{ci} + \sum_{i=1}^k S_{ci} \right) + T_{\text{controller}} + T_{\text{P4OS}}$$

$$\text{DPDK} \quad 2 \left( \sum_{i=1}^{n+1} L_i + \sum_{i=1}^n S_i \right) + 2 \left( \sum_{i=1}^{m+1} L_{Si} + \sum_{i=1}^m S_{Si} \right) + T_{\text{DPDK}}$$




# Performance Comparison

In-Network Computing


End-Host Computing





Create / Remove a flow (Latency)


**P4**

$$2 \left( \sum_{i=1}^{n+1} L_i + \sum_{i=1}^n S_i \right) + 3 \left( \sum_{i=1}^{k+1} L_{ci} + \sum_{i=1}^k S_{ci} \right) + T_{controller} + T_{P4OS}$$


**DPDK**

$$2 \left( \sum_{i=1}^{n+1} L_i + \sum_{i=1}^n S_i \right) + 2 \left( \sum_{i=1}^{m+1} L_{Si} + \sum_{i=1}^m S_{Si} \right) + T_{DPDK}$$




# Performance Comparison

In-Network Computing


End-Host Computing




Create / Remove a flow (Latency)


**P4**


$$2 \left( \sum_{i=1}^{n+1} L_i + \sum_{i=1}^n S_i \right) + 3 \left( \sum_{i=1}^{k+1} L_{ci} + \sum_{i=1}^k S_{ci} \right) + T_{controller} + T_{P4OS}$$


**DPDK**

$$2 \left( \sum_{i=1}^{n+1} L_i + \sum_{i=1}^n S_i \right) + 2 \left( \sum_{i=1}^{m+1} L_{Si} + \sum_{i=1}^m S_{Si} \right) + T_{DPDK}$$



Process a packet (Latency)




# Performance Comparison

In-Network Computing


End-Host Computing




## Create / Remove a flow (Latency)


**P4**

$$2 \left( \sum_{i=1}^{n+1} L_i + \sum_{i=1}^n S_i \right) + 3 \left( \sum_{i=1}^{k+1} L_{ci} + \sum_{i=1}^k S_{ci} \right) + T_{controller} + T_{P4OS}$$


**DPDK**


$$2 \left( \sum_{i=1}^{n+1} L_i + \sum_{i=1}^n S_i \right) + 2 \left( \sum_{i=1}^{m+1} L_{Si} + \sum_{i=1}^m S_{Si} \right) + T_{DPDK}$$



## Process a packet (Latency)

**P4**

$$\left( \sum_{i=1}^{n+1} L_i + \sum_{i=1}^n S_i \right) + T_{P4}$$






# Performance Comparison

In-Network Computing


End-Host Computing




## Create / Remove a flow (Latency)


**P4**

$$2 \left( \sum_{i=1}^{n+1} L_i + \sum_{i=1}^n S_i \right) + 3 \left( \sum_{i=1}^{k+1} L_{ci} + \sum_{i=1}^k S_{ci} \right) + T_{controller} + T_{P4OS}$$


**DPDK**


$$2 \left( \sum_{i=1}^{n+1} L_i + \sum_{i=1}^n S_i \right) + 2 \left( \sum_{i=1}^{m+1} L_{Si} + \sum_{i=1}^m S_{Si} \right) + T_{DPDK}$$



## Process a packet (Latency)

**P4**

$$2 \left( \sum_{i=1}^{n+1} L_i + \sum_{i=1}^n S_i \right) + T_{P4}$$




# Performance Comparison

In-Network Computing


End-Host Computing




## Create / Remove a flow (Latency)


**P4**

$$2 \left( \sum_{i=1}^{n+1} L_i + \sum_{i=1}^n S_i \right) + 3 \left( \sum_{i=1}^{k+1} L_{ci} + \sum_{i=1}^k S_{ci} \right) + T_{controller} + T_{P4OS}$$


**DPDK**


$$2 \left( \sum_{i=1}^{n+1} L_i + \sum_{i=1}^n S_i \right) + 2 \left( \sum_{i=1}^{m+1} L_{Si} + \sum_{i=1}^m S_{Si} \right) + T_{DPDK}$$



## Process a packet (Latency)

**P4**

$$2 \left( \sum_{i=1}^{n+1} L_i + \sum_{i=1}^n S_i \right) + T_{P4}$$



DPDK

# Performance Comparison

In-Network Computing


End-Host Computing



## Create / Remove a flow (Latency)

**P4**  $2 \left( \sum_{i=1}^{n+1} L_i + \sum_{i=1}^n S_i \right) + 3 \left( \sum_{i=1}^{k+1} L_{ci} + \sum_{i=1}^k S_{ci} \right) + T_{controller} + T_{P4OS}$


**DPDK**  $2 \left( \sum_{i=1}^{n+1} L_i + \sum_{i=1}^n S_i \right) + 2 \left( \sum_{i=1}^{m+1} L_{Si} + \sum_{i=1}^m S_{Si} \right) + T_{DPDK}$



## Process a packet (Latency)

**P4**  $2 \left( \sum_{i=1}^{n+1} L_i + \sum_{i=1}^n S_i \right) + T_{P4}$

**DPDK**  $2 \left( \sum_{i=1}^{n+1} L_i + \sum_{i=1}^n S_i \right) + 2 \left( \sum_{i=1}^{m+1} L_{Si} + \sum_{i=1}^m S_{Si} \right) + T_{DPDK}$




# Performance Comparison

In-Network Computing


End-Host Computing




## Create / Remove a flow (Latency)


**P4**


$$2 \left( \sum_{i=1}^{n+1} L_i + \sum_{i=1}^n S_i \right) + 3 \left( \sum_{i=1}^{k+1} L_{ci} + \sum_{i=1}^k S_{ci} \right) + T_{controller} + T_{P4OS}$$


**DPDK**


$$2 \left( \sum_{i=1}^{n+1} L_i + \sum_{i=1}^n S_i \right) + 2 \left( \sum_{i=1}^{m+1} L_{Si} + \sum_{i=1}^m S_{Si} \right) + T_{DPDK}$$




## Process a packet (Latency)


**P4**

$$2 \left( \sum_{i=1}^{n+1} L_i + \sum_{i=1}^n S_i \right) + T_{P4}$$


**DPDK**

$$2 \left( \sum_{i=1}^{n+1} L_i + \sum_{i=1}^n S_i \right) + 2 \left( \sum_{i=1}^{m+1} L_{Si} + \sum_{i=1}^m S_{Si} \right) + T_{DPDK}$$



# Conclusion

# Conclusion



**Elephant Flows : in-network computing**



university of  
 groningen

faculty of science  
 and engineering

bernoulli institute

# Conclusion



**Elephant Flows : in-network computing**



**Mice Flows : end-host computing**



university of  
 groningen

faculty of science  
 and engineering

bernoulli institute

# Conclusion



**Elephant Flows : in-network computing**



**Mice Flows : end-host computing**

	P4	$\Leftrightarrow$ DPDK
Performance	✓	⚠
Flexibility	⚠	✓



# Conclusion



**Elephant Flows : in-network computing**



**Mice Flows : end-host computing**

**P4 Constraints :**

- **Lack of loops**
- **Lack of support for floats**

**P4-Target Constraints :**

- **Limited memory**
- **Lack of support for non-trivial operations, such as**
  - **Multiplication**
  - **Division**
  - **Arbitrary bit shifting**

	P4	$\Leftrightarrow$ DPDK
Performance	✓	⚠
Flexibility	⚠	✓

# Conclusion



**Elephant Flows : in-network computing**



**Mice Flows : end-host computing**

## P4 Constraints :

- **Lack of loops**
- **Lack of support for floats**


## P4-Target Constraints :

- **Limited memory**
- **Lack of support for non-trivial operations, such as**
  - **Multiplication**
  - **Division**
  - **Arbitrary bit shifting**

	P4	↔	DPDK
Performance	✓		⚠
Flexibility	⚠		✓

## Viability depends on :

- **The ability of the operator to conform to the P4 and P4-target constraints**
- **The topology of the network**
- **The size of the flows**



university of  
groningen

faculty of science  
and engineering

bernoulli institute